

Laxmi Charitable Trust's
Sheth L.U.J College of Arts & Sir M.V. College of Science and Commerce
Department of Information Technology (B.Sc.I.T Semester V)
Applied Business Analytics

Practical 3

Roll No.: T018	Name: Radhika Mali
Class: TYIT	Batch: 1
Date of Assignment: 14/08/26	Date/Time of Submission: 14/08/26

AIM: Finding Relationships Among Variables:

- A. Use a dataset with multiple variables and perform correlation analysis to determine the strength and direction of relationships between pairs of variables.

CODE:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load Dataset

df = pd.read_csv("housing.csv")

# Display First Five Records

print(df.head())

# Display Numerical Columns

print(df.select_dtypes(include="number").columns)
```

```
***
   longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0    -122.23    37.88             41.0         880.0         129.0
1    -122.22    37.86             21.0       7099.0        1106.0
2    -122.24    37.85             52.0        1467.0         190.0
3    -122.25    37.85             52.0        1274.0         235.0
4    -122.25    37.85             52.0        1627.0         280.0

   population  households  median_income  median_house_value  ocean_proximity
0         322.0         126.0         8.3252        452600.0      NEAR BAY
1        2401.0        1138.0         8.3014        358500.0      NEAR BAY
2         496.0         177.0         7.2574        352100.0      NEAR BAY
3         558.0         219.0         5.6431        341300.0      NEAR BAY
4         565.0         259.0         3.8462        342200.0      NEAR BAY
Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
      'total_bedrooms', 'population', 'households', 'median_income',
      'median_house_value'],
      dtype='object')
```

```
***
   longitude  longitude  latitude  housing_median_age  total_rooms  \
longitude      1.000000  -0.924664      -0.108197      0.044568
latitude      -0.924664      1.000000      0.011173      -0.036100
housing_median_age  -0.108197      0.011173      1.000000      -0.361262
total_rooms        0.044568      -0.036100      -0.361262      1.000000
total_bedrooms     0.069608      -0.066983      -0.320451      0.930380
population         0.099773      -0.108785      -0.296244      0.857126
households         0.055310      -0.071035      -0.302916      0.918484
median_income      -0.015176      -0.079809      -0.119034      0.198050
median_house_value  -0.045967      -0.144160      0.105623      0.134153

   total_bedrooms  population  households  median_income  \
longitude      -0.069608      0.099773      0.055310      -0.015176
latitude      -0.066983      -0.108785      -0.071035      -0.079809
housing_median_age  -0.320451      -0.296244      -0.302916      -0.119034
total_rooms      0.930380      0.857126      0.918484      0.198050
total_bedrooms    1.000000      0.877747      0.979728      -0.007723
population       0.877747      1.000000      0.907222      0.004834
households       0.979728      0.907222      1.000000      0.013033
median_income    -0.007723      0.004834      0.013033      1.000000
median_house_value  0.045967      -0.024650      0.065843      0.688075

   median_house_value
longitude      -0.045967
latitude      -0.144160
housing_median_age  0.105623
total_rooms     0.134153
total_bedrooms   0.049686
population      -0.024650
households       0.065843
```

```
# Calculate Correlation Matrix

correlation = df.corr(numeric_only=True)
print(correlation)
```

OUTPUT:


```
# Display Correlation Heatmap

plt.figure(figsize=(10, 7))

sns.heatmap(
    correlation,
    annot=True,
    cmap="coolwarm"
)

plt.title("Correlation Heatmap")

plt.show()
```



- B. Apply regression analysis to identify the relationship between an independent variable (e.g., advertising expenditure) and a dependent variable (e.g., sales revenue).

Code:

```
# Import Libraries

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Load Dataset

df = pd.read_csv("housing.csv")

# Select Independent and Dependent Variables

X = df[["median_income"]]
y = df["median_house_value"]
```

```
# Split Dataset into Training and Testing Data

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42
)
```

```
# Create Linear Regression Model

model = LinearRegression()
```

```
# Train the Model

model.fit(X_train, y_train)

LinearRegression
```

```
# Display Coefficient and Intercept

print("Coefficient:", model.coef_[0])
print("Intercept:", model.intercept_)
```

```
Coefficient: 41933.84939381272
Intercept: 44459.72916907875
```

```
# Predict House Values
y_pred = model.predict(X_test)
```

```
# Display First 10 Predictions
print("Predicted Values:")
print(y_pred[:10])
```

```
Predicted Values:
[114958.91676996 150606.88213964 190393.71844449 285059.38345102
 200663.31816103 242165.24890609 257647.22610228 199229.18051176
 245893.1681172  384677.43607096]
```

```
# Calculate Mean Squared Error
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

```
Mean Squared Error: 7091157771.76555
```

```
# Calculate R2 Score
r2 = r2_score(y_test, y_pred)
print("R2 Score:", r2)
```


```
R2 Score: 0.45885918903846656
```

Output:

```
# Scatter Plot with Regression Line
plt.figure(figsize=(8, 5))

sns.regplot(
    x="median_income",
    y="median_house_value",
    data=df
)

plt.xlabel("Median Income")
plt.ylabel("Median House Value")
plt.title("Median Income vs Median House Value")
plt.show()
```



Additional Questions on housing.csv (Must include in Print document)

Q.1 Calculate the complete correlation matrix for all numerical variables.

Identify:

The strongest positive correlation.

The strongest negative correlation.

The weakest correlation.

Display the results.

Code:

Output:

```

import pandas as pd
import numpy as np
df = pd.read_csv("housing - housing.csv")
# Select numerical variables
numerical_data = df.select_dtypes(include=np.number)
# Calculate correlation matrix
correlation_matrix = numerical_data.corr()
print("Complete Correlation Matrix:")
print(correlation_matrix)
# Get correlation pairs
pairs = []
columns = correlation_matrix.columns
for i in range(len(columns)):
    for j in range(i + 1, len(columns)):
        pairs.append(
            (columns[i], columns[j], correlation_matrix.iloc[i, j])
        )
# Strongest positive correlation
strongest_positive = max(pairs, key=lambda x: x[2])
# Strongest negative correlation
strongest_negative = min(pairs, key=lambda x: x[2])
# Weakest correlation
weakest = min(pairs, key=lambda x: abs(x[2]))
print("\nStrongest Positive Correlation:")
print(strongest_positive)
print("\nStrongest Negative Correlation:")
print(strongest_negative)
print("\nWeakest Correlation:")
print(weakest)

```

Complete Correlation Matrix:

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.000000	-0.924664	-0.108197	0.044568	
latitude	-0.924664	1.000000	0.011173	-0.036100	
housing_median_age	-0.108197	0.011173	1.000000	-0.361262	
total_rooms	0.044568	-0.036100	-0.361262	1.000000	
total_bedrooms	0.069608	-0.066983	-0.320451	0.930380	
population	0.099773	-0.108785	-0.296244	0.857126	
households	0.055310	-0.071035	-0.302916	0.918484	
median_income	-0.015176	-0.079809	-0.119034	0.198050	
median_house_value	-0.045967	-0.144160	0.105623	0.134153	

	total_bedrooms	population	households	median_income	\
longitude	0.069608	0.099773	0.055310	-0.015176	
latitude	-0.066983	-0.108785	-0.071035	-0.079809	
housing_median_age	-0.320451	-0.296244	-0.302916	-0.119034	
total_rooms	0.930380	0.857126	0.918484	0.198050	
total_bedrooms	1.000000	0.877747	0.979728	-0.007723	
population	0.877747	1.000000	0.907222	0.004834	
households	0.979728	0.907222	1.000000	0.013033	
median_income	-0.007723	0.004834	0.013033	1.000000	
median_house_value	0.049686	-0.024650	0.065843	0.688075	

	median_house_value
longitude	-0.045967
...	
('longitude', 'latitude', np.float64(-0.924664339150368))	

Q.2 Build a regression model using:

X = median_income

X = housing_median_age

X = total_rooms

X = population

X = households

Y = median_house_value

Evaluate the model using:

- R2
- Mean Squared Error

Compare this model with the simple regression model using only median_income.

Code:

Output:

```

import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error
df = pd.read_csv("housing - housing.csv")
# Independent variables
X = df[
    [
        "median_income",
        "housing_median_age",
        "total_rooms",
        "population",
        "households"
    ]
]
# Dependent variable
y = df["median_house_value"]
# Multiple Linear Regression
model = LinearRegression()
model.fit(X, y)
# Predictions
y_pred = model.predict(X)
# Evaluation
r2 = r2_score(y, y_pred)
mse = mean_squared_error(y, y_pred)
print("Multiple Linear Regression")
print("-----")
print("Intercept:", model.intercept_)
print("\nCoefficients:")
for feature, coefficient in zip(X.columns, model.coef_):
    print(feature, ":", coefficient)
print("\nR2 Score:", r2)
print("Mean Squared Error:", mse)

```

Multiple Linear Regression

Intercept: -37926.606155747344

Coefficients:

median_income : 45969.659381936
 housing_median_age : 1842.1494875739568
 total_rooms : -13.946204039687927
 population : -39.44546486713395
 households : 214.35386822082083

R2 Score: 0.5629054644856253

Mean Squared Error: 5820133599.282842

Q.3 Find the correlation between:

- total_rooms
- median_house_value

Code:

```

# Simple Linear Regression
X_simple = df[["median_income"]]
simple_model = LinearRegression()
simple_model.fit(X_simple, y)
# Predictions
y_pred_simple = simple_model.predict(X_simple)
# Evaluation
r2_simple = r2_score(y, y_pred_simple)
mse_simple = mean_squared_error(y, y_pred_simple)
print("Simple Linear Regression")
print("-----")
print("R2 Score:", r2_simple)
print("Mean Squared Error:", mse_simple)

```

Output:

Simple Linear Regression

R2 Score: 0.47344749180719903

Mean Squared Error: 7011311502.929527